

# NoSQL

**Benmounah Zakaria**  
**Constantine 2 University**  
**SDIA M1**  
**2023-2024**

---

# **PART 01**

## **HISTORY ERA BEFORE NOSQL**

# Pre-2002 Era

- Most organizations had their own relational databases.
- They would buy expensive servers (e.g., \$40,000 computers) to host and manage their databases. This was referred to as vertical scaling.
- Some vendors claimed to be "cloud" by merely collecting individual databases and hosting them on platforms like Amazon. This wasn't true cloud computing but more of a collection of individual databases.

# Web Search and Crawling

**Late 1990s - Early 2000s**



- Google set out to index the entire internet.
- Traditional databases couldn't handle such a vast amount of data.
- Google's solution was to crawl the web, make a full copy, and then index that copy.
- This endeavor required new data management and storage techniques, as the scale was unlike anything seen before in the database world.

# Gmail Introduction

April 1, 2004



- When Google introduced Gmail, it wasn't just a new email service; it was a challenge to the established norms of data management.
- Traditional relational databases were designed for smaller, isolated workloads.
- In contrast, Gmail had to serve millions of users globally, providing each with a significant amount of storage (1 GB at launch, which was unprecedented at the time).
- Gmail's architecture needed to be highly distributed and scalable. Each user's data had to be isolated, yet the system had to be responsive and reliable.

# Data Silos and Eventual Consistency

## 2000s onwards



- Google's applications, including Gmail, were designed to operate with **data silos**.
- Each user's interactions were restricted to their own "silo" or segment of data.
- For instance, if a user deleted an email, it affected only their silo and not others.
- This allowed for "eventual consistency," where data didn't need to be immediately consistent across all servers but would eventually reach a consistent state.
- This approach was different from traditional databases, which prioritized immediate consistency (ACID properties).

# Dynamic Data Migration

## 2000s onwards



- In the early days of Gmail, Google had the capability to migrate a user's data closer to their geographical location.
- If a user traveled from the US to Europe, their data might shift to European servers, making access faster.
- This dynamic migration was innovative and showcased Google's commitment to user experience.
- Google's approach to handling large-scale, user-centric applications required rethinking traditional database architectures.

# Revelations from Google



- In conferences like Google I/O 2008, Google began to share insights into how they built their systems. These revelations were revolutionary for many in the tech industry.
- The speaker was particularly impressed by Google's gather/scatter technique, where Google used numerous cheaper computers to achieve faster performance than one could get from a single expensive machine.
- By 2010, Google started showcasing more of their techniques, emphasizing efficiency and sustainability. They recognized the importance of sharing best practices for the greater good.



# Observe where the story first started...



- <https://www.youtube.com/watch?v=Md7K90FfJhg>
- <https://www.youtube.com/watch?v=Ho1GEyftpmQ>
- <https://www.youtube.com/watch?app=desktop&v=6x0cAzQ7PVs>

# Marissa Mayer Video



- Google search's perceived intelligence operates within a timeframe that's indistinguishable for users,
- whether it takes a tenth of a second or a thousandth of a second.
- The scatter-gather approach is efficient even if it takes a quarter of a second, as the resources used during that time are minuscule.
- This approach ensures consistency in search speed no matter the search volume.

# Container Tour Video



- Google's data centers are built for efficiency, not aesthetics.
- There's no unnecessary packaging or wasted energy.
- The computers are designed to be disposable and modular, emphasizing replication to ensure data security.
- Each user's Gmail data, for example, is replicated multiple times across various servers and locations.

# Matt Cutts Video




- Google's approach to computing **page rank** is elegantly simple when understood.
- It is a distributed algorithm, designed to converge and adjust as new pages emerge.
- Google's approach is unique, opting for distributed small databases rather than one large centralized one.

# Shift to Amazon & AWS



- Amazon's evolution from a book company to a tech giant with AWS changed the tech landscape.
- Instead of investing in expensive hardware that becomes obsolete, the idea of renting space and computational power on cheap, commodity hardware became appealing.
- This led to the rise of "carpet clusters" – using networked commodity computers to create a scalable, distributed system.

# Second generation



- The transition from the first generation, exemplified by Gmail and Google Search, to the second, represented by Facebook and Twitter.
- Gmail and Google Search operated largely in user-specific silos, with interactions relatively simple.
- In contrast, platforms like Facebook and Twitter, which introduced friend networks and followers, necessitated a more interconnected and dynamically updated database.

# Challenges with Facebook



- Unlike Gmail's isolated data silos, Facebook required constant data migration and replication due to its interactivity.
- For instance, when a user likes a post, this "like" has to be reflected in real-time across various servers catering to different friends.
- This brings up a myriad of questions and challenges, especially when factoring in privacy settings and friend connections.

# Eventual Consistency



- Facebook is an example of an "eventual consistency" database.
- This means that while a user might like a post, it might take some time for this action to reflect across all relevant servers.
- Such databases prioritize overall system health over real-time accuracy.



# Engineer's Dilemma

- The engineers often overgeneralize solutions.
- Once they've crafted a solution for a specific problem, they might try to adapt it for broader applications, leading to the creation of new database forms.
- The engineer, when asked to pass the salt, begins developing a system to pass any condiment, believing that it'll be more efficient in the future.
- While the engineer's intention might be to create a versatile solution, it's an overcomplicated response to a simple request.

# Emergence of New Database Forms

- The success of second-generation cloud companies led to the development of **BASE-style databases**, moving away from traditional ACID databases.

# **PART 02**

## **BASE-style**

# BASE-style databases

- **Basically Available:** The system ensures availability, even in the presence of multiple failures.
- **Soft state:** The state of the system may change over time, even without input. This reflects the inherent uncertainty and inconsistency in distributed systems.
- **Eventual Consistency:** The system will become consistent over time, given that the system doesn't receive input during that time.

# BASE-style vs ACID-style

- **ACID** (Atomicity, Consistency, Isolation, Durability) databases, like Oracle, Postgres, **MySQL**, **SQLite**, and **SQLServer**, ensure that data remains consistent over time.
  - They operate on the principle that data, once written, remains the same for all readers at any given moment.
- **BASE** (Basically Available, Soft state, Eventually consistent) databases, like **Mongo**, **Cassandra**, and **Google's BigTable**, operate on eventual consistency.
  - Data might be inconsistent across nodes for a while, but it will eventually synchronize.

# BASE-style databases characterized

- Wide distribution without central locks.
- Use of fast networks, low-memory CPUs, and a multitude of disk drives.
- Data sharding with indexes for data location.
- Preference for documents over rows and columns.
- Flexible schemas allowing on-the-fly addition of key-value pairs.
- Emphasis on schema-on-read rather than rigid early schemas.

# JSON's Role

- **JSON** became a crucial format, offering simple representation of key-value pairs.
- Provides fast parsing across various languages.
- Highly compressible, making it attractive for databases aiming for memory efficiency.

# JSON example

```
{  
  "user": {  
    "id": 1,  
    "name": "John Doe",  
    "email": "john.doe@example.com"  
  }  
}
```

- This JSON structure represents a user with an id, name, and email.



# Emergence of NoSQL Databases

- **CouchDB:** Aims for a cluster of unreliable commodity hardware.
- **MongoDB:** Popularized with the Node ecosystem, using JSON storage.
- **Cassandra:** Developed by Facebook engineers based on their experiences.
- **Elasticsearch:** Evolved from an effort to replicate Google search, focusing more on indexing than being a typical database.

# Software as a Service (SaaS) Trend

- Offerings like Amazon's DynamoDB, Google's BigTable, and Microsoft's Azure allowed companies to use databases without heavy investments in in-house expertise or hardware.

# Adoption by Major Players

- **Amazon:** uses Amazon Redshift.
- **Facebook:** uses Cassandra and RocksDB.
- **Twitter:** migrated from MySQL to Manhattan, its real-time, multi-tenant distributed database.
- **Apple:** Apple replaced some of its use of Cassandra with its in-house designed FoundationDB.
- **Netflix:** heavily uses Cassandra for its scalability and distributed architecture.
- **LinkedIn:** uses kafka and Espresso.
- **Uber:** Uber uses a mix of databases, including MySQL, Postgres, and Cassandra,
- **Microsoft:** uses Azure Cosmos DB.

# PART 03

## NoSQL

# NoSQL

- NoSQL (often interpreted as “not only SQL”) databases are non-relational databases that provide a way to store and retrieve data differently than traditional relational databases.
- The term “NoSQL” is a bit of a misnomer since many of these databases do use some form of SQL-like query language.
- The primary distinction is their non-relational architecture.
- There are several types of NoSQL databases: Document-based, Column-based databases, Key-value stores, and Graph databases.