



Lecture 2

B1 Optimization - Trinity 2019 / Michaelmas 2020 - V. A. Prisacariu

- Rosenbrock's function
- Newton's method
 - Line search
- Quasi-Newton methods
- Least-Squares and Gauss-Newton methods and a few words on Levenberg-Marquardt.
- Downhill simplex (amoeba) algorithm

Rosenbrock's function

$$f(x, y) = 100(y - x^2)^2 + (1 - x)^2$$



Minimum is at [1, 1]

An Optimization Algorithm (from lecture 1)

Start at \mathbf{x}_0 then repeat:

1. Compute a search direction \mathbf{p}_k .
2. Compute a step length α_k , such that $f(\mathbf{x}_k + \alpha_k \mathbf{p}_k) < f(\mathbf{x}_k)$.
3. Update $\mathbf{x}_{k+1} = \mathbf{x}_k + \alpha_k \mathbf{p}_k$
4. Check for convergence (termination criteria), e.g. $\nabla f \approx 0$.

Reduce optimization in N dimensions to a series of (1D) line minimizations.

Steepest descent (from lecture 1)

Basic principle is to minimize the N-dimensional function by a series of 1D line-minimizations:



$$\mathbf{x}_{n+1} = \mathbf{x}_n + \alpha_n \mathbf{p}_n$$

The steepest decent method choses \mathbf{p}_n to be parallel to the negative gradient:

$$\mathbf{p}_n = -\nabla f(\mathbf{x}_n)$$

Steepest descent

The 1D line minimization must be performed using one of the earlier methods



- The zig-zag behavior is clear in the zoomed view (100 iterations);
- The algorithm crawls along the valley.

Newton's method in 1D (from lecture 1)

Fit a quadratic approx. to $f(x)$ using both gradient and curvature information at x .

- Expand $f(x)$ locally using a Taylor series:

$$f(x + \delta x) = f(x) + \delta x f'(x) + \frac{\delta x^2}{2} f''(x) + \text{h.o.t.}$$

- Find the δx such that $x + \delta x$ is a stationary point of f :

$$\frac{d}{d\delta x} \left(f(x) + \delta x f'(x) + \frac{\delta x^2}{2} f''(x) \right) = f'(x) + \delta x f''(x) = 0$$

- and rearranging:

$$\delta x = -\frac{f'(x)}{f''(x)}$$

- Update for x :

$$x_{n+1} = x_n - \frac{f'(x_n)}{f''(x_n)}$$

Taylor expansion in 2D (from lecture 1)

A function can be approximated locally by its Taylor series expansion about a point x_0 .

$$f(x_0 + x) \approx f(x_0) + \left(\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right) \begin{pmatrix} x \\ y \end{pmatrix} + \frac{1}{2} (x, y) \begin{bmatrix} \frac{\partial^2 f}{\partial x^2} & \frac{\partial^2 f}{\partial x \partial y} \\ \frac{\partial^2 f}{\partial x \partial y} & \frac{\partial^2 f}{\partial y^2} \end{bmatrix} \begin{pmatrix} x \\ y \end{pmatrix} + \text{h.o.t.}$$

- This is a generalization of the 1D Taylor series:

$$f(x + \delta x) = f(x) + \delta x f'(x) + \frac{\delta x^2}{2} f''(x) + \text{h.o.t.}$$

- The expansion to second order is a **quadratic** function in \mathbf{x} .

$$f(\mathbf{x}) = a + \mathbf{g}^T \mathbf{x} + \frac{1}{2} \mathbf{x}^T \mathbf{H} \mathbf{x}$$

Newton's method in ND

Expand $f(\mathbf{x})$ by its Taylor series about the point \mathbf{x}_n

$$f(\mathbf{x}_n + \delta\mathbf{x}) \approx f(\mathbf{x}_n) + \mathbf{g}_n^T \delta\mathbf{x} + \frac{1}{2} \delta\mathbf{x}^T \mathbf{H}_n \delta\mathbf{x}$$

where the gradient is the vector

$$\mathbf{g}_n = \nabla f(\mathbf{x}) = \left[\frac{\partial f}{\partial x_1}, \dots, \frac{\partial f}{\partial x_N} \right]^T$$

and the Hessian is the symmetric matrix

$$\mathbf{H}_n = \mathbf{H}(\mathbf{x}_n) = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \dots & \frac{\partial^2 f}{\partial x_1 \partial x_N} \\ \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_1 \partial x_N} & \dots & \frac{\partial^2 f}{\partial x_N^2} \end{bmatrix}$$

For a minimum we will require that $\nabla f(\mathbf{x}) = 0$, and so

$$\nabla f(\mathbf{x}) = \mathbf{g}_n + \mathbf{H}_n \delta\mathbf{x} = 0$$

with the solution $\delta\mathbf{x} = -\mathbf{H}_n^{-1} \mathbf{g}_n$. This gives the iterative update

$$\mathbf{x}_{n+1} = \mathbf{x}_n - \mathbf{H}_n^{-1} \mathbf{g}_n$$

Newton's method in ND



$$\mathbf{x}_{n+1} = \mathbf{x}_n + \delta\mathbf{x} = \mathbf{x}_n - \mathbf{H}_n^{-1}\mathbf{g}_n$$

- If $f(\mathbf{x})$ is quadratic, the solution is found in one step.
- The method has quadratic convergence (as in the 1D case).
- The solution $\delta\mathbf{x} = -\mathbf{H}_n^{-1}\mathbf{g}_n$ is guaranteed to be a downhill direction (provided that \mathbf{H} is positive definite).
- For numerical reasons, the inverse is not actually computed, instead $\delta\mathbf{x}$ is computed as the solution of $\mathbf{H}\delta\mathbf{x} = -\mathbf{g}_n$.
- Rather than jump straight to $\mathbf{x}_n - \mathbf{H}_n^{-1}\mathbf{g}_n$, it is often better to perform a line search which ensures (more) global convergence:

$$\mathbf{x}_{n+1} = \mathbf{x}_n - \alpha\mathbf{H}_n^{-1}\mathbf{g}_n$$

- If $\mathbf{H} = \mathbf{I}$, Newton reduces to steepest descent.

Newton's method in ND Example




ellipses show successive
quadratic approximations

- The algorithm converges in only 15 iterations compared to hundreds for steepest descent.
- **However**, the method requires computing the Hessian matrix at each iteration – this is not always feasible

Quasi-Newton Methods

- If the problem size is large and the Hessian matrix is dense, then it may be infeasible / inconvenient to compute it directly.
- Quasi-Newton method avoid this problem by keeping “rolling estimate” of $H(\mathbf{x})$, updated at each iteration using new gradient information.
- Common schemes are Broyden–Fletcher–Goldfarb–Shanno (**BFGS**) and Davidon–Fletcher–Powell (**DFP**).



First derivatives:

$$f' \left(x_0 + \frac{h}{2} \right) = \frac{f_1 - f_0}{h} \text{ and } f' \left(x_0 - \frac{h}{2} \right) = \frac{f_0 - f_{-1}}{h}$$

Second derivatives:

$$f''(x_0) = \frac{\frac{f_1 - f_0}{h} - \frac{f_0 - f_{-1}}{h}}{h} = \frac{f_1 - 2f_0 + f_{-1}}{h^2}$$

For H_{n+1} we can build an approximation from $H_n, \mathbf{g}_n, \mathbf{g}_{n+1}, \mathbf{x}_n, \mathbf{x}_{n+1}$

Quasi-Newton BFGS

Set $H_0 = I$.

Update according to

$$H_{n+1} = H_n + \frac{\mathbf{q}_n \mathbf{q}_n^T}{\mathbf{q}_n^T \mathbf{s}_n} - \frac{(H_n \mathbf{s}_n)(H_n \mathbf{s}_n)^T}{\mathbf{s}_n^T H_n \mathbf{s}_n}$$


where

$$\mathbf{s}_n = \mathbf{x}_{n+1} - \mathbf{x}_n$$

$$\mathbf{q}_n = \mathbf{g}_{n+1} - \mathbf{g}_n$$

- The matrix itself is not stored, but rather represented compactly by a few stored vectors.
- The estimate H_{n+1} is used to form a local quadratic approximation as before.

Quasi-Newton BFGS Example



- The method converges in 25 iterations, compared to 15 for the full-Newton method
- In Matlab the optimization function '**fminunc**' uses a BFGS quasi-Newton method for medium scale optimization problems.

Matlab – fminunc

```
>> f='100*(x(2)-x(1)^2)^2+(1-x(1))^2';
```

```
>> GRAD='[100*(4*x(1)^3-4*x(1)*x(2))+2*x(1)-2; 100*(2*x(2)-2*x(1)^2) ]';
```

Choose options for BFGS quasi-Newton

```
>> OPTIONS=optimset('LargeScale','off', 'HessUpdate','bfgs' );
```

```
>> OPTIONS = optimset(OPTIONS,'gradobj','on');
```

Start point

```
>> x = [-1.9; 2];
```

```
>> [x,fval] = fminunc({f,GRAD},x,OPTIONS);
```

This produces

```
x = 0.9998, 0.9996      fval = 3.4306e-008
```

Non-Linear Least Squares

- It is **very** common in applications for a cost function $f(\mathbf{x})$ to be the sum of a large number of squared residuals:

$$f(x) = \sum_{i=1}^M r_i^2$$

- If each residual depends **non-linearly** on the parameters \mathbf{x} , then the minimization of $f(\mathbf{x})$ is a non-linear least squares problem.
- Examples arise in non-linear regression (fitting) of data.

Linear Least Squares Reminder



The goal is to fit a smooth curve to measured data points $\{s_i, t_i\}$ by minimizing the cost:

$$\begin{aligned} f(\mathbf{x}) &= \sum_{i=1} r_i^2 \\ &= \sum_{i=1} (y(s_i, \mathbf{x}) - t_i)^2, \text{ where } t_i \text{ is the target value} \end{aligned}$$

For example, the regression functions $y(s_i, \mathbf{x})$ might be the polynomial

$$y(s, \mathbf{x}) = x_0 + x_1 s + x_2 s^2 + \dots$$

In this case the function is linear in the parameter \mathbf{x} and there is a closed form solution. In general there will be no closed form solution to non-linear $y(s, \mathbf{x})$.



Non-Linear Least Squares Example: Aligning a 3D Model to an Image



Input:

3D textured face model, camera model, image $I(x, y)$.



Task:

Determine the 3D rotation and 3D translation that minimizes the error between the image $I(x, y)$ and the projected 3D model.

Cost Function


$$f(\mathbf{R}, \mathbf{T}) = \sum_{x,y}^M \left| \hat{I}_{\mathbf{R},\mathbf{T}}(x,y) - I(x,y) \right|^2 \quad N=6, M = 10^6$$


Transformation parameters:

- 3D rotation matrix \mathbf{R}
- Translation vector $\mathbf{T} = (T_x, T_y, T_z)^T$

Image generation:

- Rotate and translate 3D model by \mathbf{R} and \mathbf{T}
- Project to generate image $\hat{I}_{\mathbf{R},\mathbf{T}}(x,y)$







Non-Linear Least Squares

$$f(\mathbf{x}) = \sum_{i=1}^M r_i^2 = |\mathbf{r}|^2$$

The $M \times N$ Jacobian of the vector of residuals \mathbf{r} is defined as:

$$\mathbf{J}(\mathbf{x}) = \begin{pmatrix} \frac{\partial r_1}{\partial x_1} & \cdots & \frac{\partial r_1}{\partial x_N} \\ \vdots & \ddots & \vdots \\ \frac{\partial r_M}{\partial x_1} & \cdots & \frac{\partial r_M}{\partial x_N} \end{pmatrix}$$

assume
 $M > N$

$$\begin{pmatrix} \mathbf{J} \end{pmatrix}$$

Consider

$$\frac{\partial}{\partial x_k} \sum_i r_i^2 = \sum_i 2r_i \frac{\partial r_i}{\partial x_k}$$

Hence

$$\nabla f(\mathbf{x}) = 2\mathbf{J}^T \mathbf{r}$$

$$\begin{pmatrix} \end{pmatrix} = \begin{pmatrix} \mathbf{J}^T \end{pmatrix} \begin{pmatrix} \end{pmatrix}$$

Non-Linear Least Squares

For the Hessian we require

$$\frac{\partial^2}{\partial x_l \partial x_k} \sum_i r_i^2 = 2 \frac{\partial}{\partial x} \sum_i r_i \frac{\partial r_i}{\partial x_k} = 2 \sum_i \frac{\partial r_i}{\partial x_k} \frac{\partial r_i}{\partial x_l} + 2 \sum_i r_i \frac{\partial^2 r_i}{\partial x_k \partial x_l}$$

Hence

$$H(\mathbf{x}) = 2\mathbf{J}^T \mathbf{J} + 2 \sum_i^M r_i \mathbf{R}_i$$
$$\begin{pmatrix} \mathbf{J}^T \end{pmatrix} \begin{pmatrix} \mathbf{J} \end{pmatrix}$$

Non-Linear Least Squares

Note that the second-order term in the Hessian $H(\mathbf{x})$ is multiplied by the residuals r_i .

In most problems, the residuals will typically be small.

Also, at the minimum, the residuals will typically be distributed with mean = 0.

For these reasons, the second-order term is often ignored, giving the **Gauss-Newton** approximation of the Hessian:

$$H(\mathbf{x}) = 2J^T J$$

Hence, explicit computation of the full Hessian can be again avoided.

Gauss-Newton Example

The minimization of the Rosenbrock function.

$$f(x, y) = 100(y - x^2)^2 + (1 - x)^2$$

can we be written as a least-squares problem with residual vector

$$\mathbf{r} = \begin{bmatrix} 10(y - x^2) \\ (1 - x) \end{bmatrix}$$

$$\mathbf{J}(\mathbf{x}) = \begin{pmatrix} \frac{\partial r_1}{\partial x} & \frac{\partial r_1}{\partial y} \\ \frac{\partial r_2}{\partial x} & \frac{\partial r_2}{\partial y} \end{pmatrix} = \begin{pmatrix} -20x & 10 \\ -1 & 0 \end{pmatrix}$$

Gauss-Newton Example

The true Hessian is:

$$H(\mathbf{x}) = \begin{bmatrix} \frac{\partial^2 f}{\partial x^2} & \frac{\partial^2 f}{\partial x \partial y} \\ \frac{\partial^2 f}{\partial x \partial y} & \frac{\partial^2 f}{\partial y^2} \end{bmatrix} = \begin{bmatrix} 1200x^2 - 400y + 2 & -400x \\ -400x & 200 \end{bmatrix}$$

The Gauss-Newton approximation of the Hessian is:

$$2J^T J = 2 \begin{bmatrix} -20x & -1 \\ 10 & - \end{bmatrix} \begin{bmatrix} -20x & 10 \\ -1 & 0 \end{bmatrix} = \begin{bmatrix} 800x^2 + 2 & -400x \\ -400x & 200 \end{bmatrix}$$

Gauss-Newton Summary

For a cost function $f(\mathbf{x})$ that is the sum of squared residuals

$$f(x) = \sum_{i=1} r_i^2$$

The Hessian can be approximated as

$$H(\mathbf{x}) = 2J^T J$$

and the gradient is given by

$$\nabla f(\mathbf{x}) = 2J^T \mathbf{r}$$

So the Newton update step

$$\mathbf{x}_{n+1} = \mathbf{x}_n + \delta\mathbf{x} = \mathbf{x}_n - H_n^{-1} \mathbf{g}_n$$



Computed as $H\delta\mathbf{x} = -\mathbf{g}_n$, becomes

$$J^T J \delta\mathbf{x} = -J^T \mathbf{r}$$

These are called the **normal equations**.

Gauss-Newton Rosenbrock


$$\mathbf{x}_{n+1} = \mathbf{x}_n - \alpha_n \mathbf{H}_n^{-1} \mathbf{g}_n \text{ with } \mathbf{H}_n(\mathbf{x}) = 2\mathbf{J}_n^T \mathbf{J}_n$$



minimization with the Gauss-Newton approximation with line search takes only 14 iterations


Comparison

Newton



- requires computing Hessian (i.e. n^2 second derivatives)
- exact solution if quadratic

Gauss-Newton



- approximates Hessian by Jacobian product
- requires only $n \times M$ first derivatives

Properties of Methods

Update: $\mathbf{x}_{n+1} = \mathbf{x}_n + \alpha_n \mathbf{p}_n$

- Gradient descent [$\mathbf{p}_n = -\mathbf{g}$]

- Will zig-zag – each new increment is perpendicular to the previous.
- Requires 1D search (unless you hack it).
- Slow to converge.

- Newton's method [$\mathbf{p}_n = -\mathbf{H}^{-1} \mathbf{g}_n$]

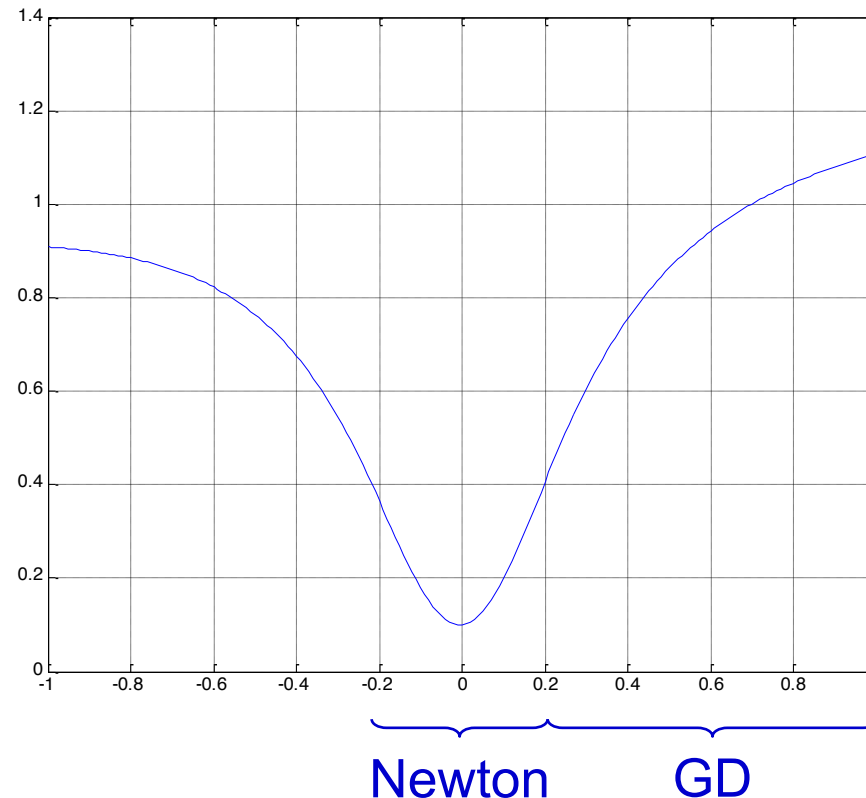
- Requires computation of Hessian.
- Can converge to maximum or saddle as well as minimum.
- Can be unstable.

- Gauss-Newton [$\mathbf{p}_n = -\mathbf{H}_{GN}^{-1} \mathbf{g}_n$]

- Is a downhill method, so will not converge to maximum or saddle.
- Can be unstable, usually needs line search.

Levenberg-Marquardt [more in C25]

- Away from the minimum, in regions of negative curvature, the Gauss-Newton approximation is not very good.
- In such regions, a simple steepest-descent step is probably the best plan.
- The **Levenberg-Marquardt** method is a mechanism for varying between steepest-descent and Gauss-Newton steps depending on how good the H_{GN} approximation is locally.





Performance issues for optimization algorithms

1. Number of iterations required
2. Cost per iteration
3. Memory footprint
4. Region of convergence

The downhill simplex (amoeba) algorithm

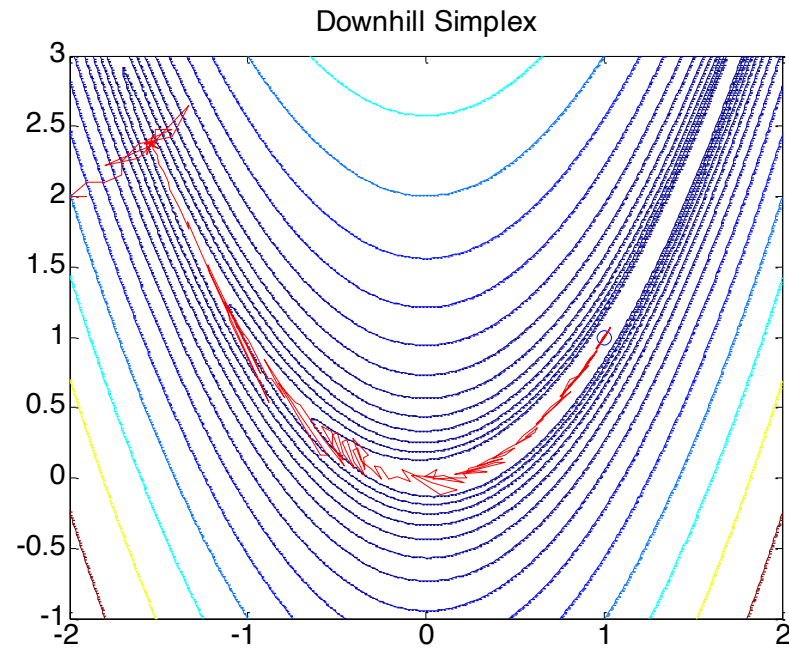
- Due to Nelder and Mead (1965)
- A *direct* method: only uses function evaluations (no derivatives)
- A simplex is the polytope in N dimensions with $N+1$ vertices, e.g.
 - 2D: triangle
 - 3D: tetrahedron
- Basic idea: move by *reflections*, *expansions* or *contractions*

The downhill simplex (amoeba) algorithm




- Reorder the points so that $f(\mathbf{x}_{n+1}) > \dots > f(\mathbf{x}_2) > f(\mathbf{x}_1)$ (i.e. \mathbf{x}_{n+1} is the worst point).
- Calculate $\bar{\mathbf{x}} = (\sum_i \mathbf{x}_i)/N$ as the centroid of all points except \mathbf{x}_{n+1} .
- Generate a trial point by **reflection** $\mathbf{x}_r = \bar{\mathbf{x}} + \alpha(\bar{\mathbf{x}} - \mathbf{x}_{n+1})$, where $\alpha > 0$.
- If the reflected point \mathbf{x}_r is neither the new best or worst point, i.e. if $f(\mathbf{x}_1) < f(\mathbf{x}_r) < f(\mathbf{x}_n)$ replace the worst point \mathbf{x}_{n+1} by \mathbf{x}_r .
- If the reflected point \mathbf{x}_r is the new best point, i.e. if $f(\mathbf{x}_r) < f(\mathbf{x}_1)$, generate a new point by **expansion** $\mathbf{x}_e = \mathbf{x}_r + \beta(\mathbf{x}_r - \bar{\mathbf{x}})$, where $\beta > 0$.
- If the expanded point is better than the reflected point $f(\mathbf{x}_e) < f(\mathbf{x}_r)$ then replace \mathbf{x}_{n+1} by \mathbf{x}_e , otherwise replace \mathbf{x}_{n+1} by \mathbf{x}_r .
- If $f(\mathbf{x}_r) > f(\mathbf{x}_n)$ assume the polytope is too large and generate a new point by **contraction** $\mathbf{x}_c = \bar{\mathbf{x}} + \gamma(\mathbf{x}_{n+1} - \bar{\mathbf{x}})$, where γ , ($0 < \gamma < 1$) is the contraction coefficient.
- If the contracted point is better than the worst point, so $f(\mathbf{x}_c) < f(\mathbf{x}_{n+1})$, the contraction has succeeded, and replace \mathbf{x}_{n+1} by \mathbf{x}_c , otherwise contract again (or shrink all points).


Downhill Simplex Example 1



Path of best vertex




Detail



Matlab fminsearch
with 200 iterations

Downhill Simplex Example 2: Contraction About a Minimum



Summary

- no derivatives required
- deals well with noise in the cost function
- is able to crawl out of some local minima (though, of course, can still get stuck)

Matlab – fminsearch

Nelder-Mead simplex direct search

```
>> banana = @(x)100*(x(2)-x(1)^2)^2+(1-x(1))^2;
```

Pass the function handle to fminsearch:

```
>> [x,fval] = fminsearch(banana,[-1.9, 2])
```

This produces

```
x = 1.0000 1.0000
```

```
fval =4.0686e-010
```

Google to find out more on using this function

What is next?

- Move from general and quadratic optimization problems to linear programming.
- Constrained optimization.